

Deep Q-Learning For The Traveling Salesman Problem

Behzad Karimi Deanta Kelly Ian Kessler Fatima Ododo

Montana State University
behzad.karimi@student.montana.edu
deanta.kelly@student.montana.edu
ian.kessler@student.montana.edu
fatima.ododo@student.montana.edu

M 508 Presentation
May 2, 2024

Introduction

- Project Focused On Implementation And Analysis Of Algorithm Described In *Learning Combinatorial Optimization Algorithms over Graphs*
- Focused On One Combinatorial Optimization (CO) Application: The Traveling Salesman Problem (TSP)
- Used Held-Karp Algorithm ($\in \mathcal{O}(2^n n^2)$) To Check Accuracy
- According to the paper, classical algorithms “seldom exploit a common trait of real-world optimization problems: instances of the same type of problem are solved again and again on a regular basis, maintaining the same combinatorial structure, but differing mainly in their data.”
- **Problem Statement:** Given a distribution \mathbb{D} of graphs, can we learn better heuristics to solve TSP that generalize to unseen instances from \mathbb{D} ?

Presentation Overview

- 1 The Traveling Salesman Problem
- 2 Q-Learning
- 3 Evaluating Q With *Structure2Vec* Neural Network
- 4 Data Analysis

The Traveling Salesman Problem

(On Whiteboard)

Q-Learning

- 1 Initialize experience replay memory M to capacity N
- 2 for episode $e = 1$ to L do
- 3 Draw graph G from distribution D
- 4 Initialize the state to empty $S_i = ()$
- 5 for step $t=1$ to T do
- 6
$$v_t = \begin{cases} \text{random node } v \in \bar{S}_t, & \text{w.p. } \epsilon \\ \operatorname{argmax}_{v \in \bar{S}_t} \hat{Q}(h(S_t), v; \theta), & \text{Otherwise} \end{cases}$$
- 7 Add v_t to partial solution: $S_{t+1} := (S_t, v_t)$
- 8 if $t \geq n$ then
- 9 Add tuple $(S_{t-n}, v_{t-n}, R_{t-n}, S_t)$ to M
- 10 Sample random batch from $B \sim M$
- 11 Update θ by SGD over (6) from B
- 12 end if
- 13 end for
- 14 end for
- 15 return θ

Structure2Vec: Describing Our Embedding

- Given partial solution S , we create a p -dimensional embedding for each node $v \in V$ and each layer i : $(\mu_S^{(i)})_v \in \mathbb{R}^{p \times 1}$
- Let $m = |V|$. Then, the i^{th} layer of our neural network is

$$\mu_S^{(i)} = [(\mu_S^{(i)})_0, (\mu_S^{(i)})_1, \dots, (\mu_S^{(i)})_{m-1}] \in \mathbb{R}^{p \times m}$$

- To have our embedding depend on S , we use

$$x_S := [1\{v \in S\} : v \in V] \in \{0, 1\}^{1 \times m}$$

- Want our embedding to exploit the graph structure, so we define
 - The Neighbors of v to be $\mathcal{N}(v) (= V \setminus \{v\}$ For Complete Graphs)
 - The Weight of Edge (v, u) to be $w(v, u)$

Structure2Vec: Calculating Hidden Layers

- Our Initial Layer: $\mu_S^{(0)} := \mathbf{0}^{p \times m}$
- From One Layer To The Next:

$$(\mu_S^{(i+1)})_v \leftarrow \text{relu}(\theta_1 x_S[v] + \theta_2 \sum_{u \in \mathcal{N}(v)} (\mu_S^{(i)})_u + \theta_3 \sum_{u \in \mathcal{N}(v)} \text{relu}(\theta_4 w(v, u)))$$

Where $\theta_1 \in \mathbb{R}^p$, $\theta_2, \theta_3 \in \mathbb{R}^{p \times p}$, $\theta_4 \in \mathbb{R}^p$

- Using T hidden layers, we compute a sequence of hidden layers:

$$\mu_S^{(0)} \rightarrow \mu_S^{(1)} \rightarrow \dots \rightarrow \mu_S^{(T)}$$

Where $\mu_S^{(T)}$ is Final Hidden Layer

Output “Q” Layer

- Using final hidden layer, $\mu_S^{(T)}$, we use evaluation function:

$$\hat{Q}(S, v) = \theta_{5a}^\top \text{relu}(\theta_6 \sum_{u \in V} (\mu_S^{(T)})_u) + \theta_{5b}^\top \text{relu}(\theta_7 (\mu_S^{(T)})_v)$$

Where $\theta_{5a}, \theta_{5b} \in \mathbb{R}^p$, $\theta_6, \theta_7 \in \mathbb{R}^{p \times p}$

- Letting $\bar{S} := V \setminus S$, we update $S \leftarrow S + [v^*]$, where

$$v^* = \operatorname{argmax}_{v \in \bar{S}} \hat{Q}(S, v)$$

- $\hat{Q}(S, v)$ depends on $\Theta = [\theta_1, \theta_2, \theta_3, \theta_4, \theta_{5a}, \theta_{5b}, \theta_6, \theta_7]$
- The parameters Θ are learned using reinforcement learning.

4-Fold Cross Validation

- Trained and Tested Model On 8 Random Euclidean Graphs

$$[G_i : i \in [0..7]]$$

- For each graph G_i , its vertex set $V_i \in \mathbb{Z} \times \mathbb{Z}$ with $|V_i| = 9$ and

$$V_i = \{(x_{ij}, y_{ij}) : -5 \leq x_{ij}, y_{ij} \leq 5\}$$

- The Error of Each Fold:
Avg. Error on $\frac{8}{4} = 2$ Graphs Using Θ Trained on $8 - 2 = 6$ Other Graphs
- Error for Model: Avg. Error of all Folds

Approximation Ratio As Error

- For a full solution S for graph G , let

$$\text{cost}_G(S) = \sum_{i=0}^{|S|-2} w(S[i], S[i+1]) + w(S[|S|-1], S[0]),$$

(The Total Weight of Tour Generated from S)

- For graph G , let ...
 - \hat{S}_Θ be a full solution found using our model with weights Θ .
 - S^* be a full solution that minimizes cost_G .
- Then, Error is the Approximation Ratio

$$\rho = \frac{\text{cost}_G(\hat{S}_\Theta)}{\text{cost}_G(S^*)} \geq 1$$

- Optimal solutions S^* were found using the Held-Karp algorithm.

Tuning Hyperparameters

- Each Deep Q-Learning Model Defined By 7 Hyperparameters
 - p : Dimension of Each Node Embedding
 - T : Number of Hidden Layers
 - ϵ : Probability of Choosing Random $v \in \overline{S}$ to Append to Partial Solution S
 - n : Number of Steps Between States For n -Step Q -Learning
 - α : Learning Rate for Gradient Descent
 - β : Maximum Size of Batches for Mini-Batch Gradient Descent
 - γ : Discount Factor For Q -Learning
- Used 2 Options For Each Hyperparameter Among Models Used
- Error Was Calculated For $2^7 = 128$ Models, Each Created From Choosing Between 2 Values For Each of The 7 Hyperparameters

Avg. Error By Hyperparameter (1)

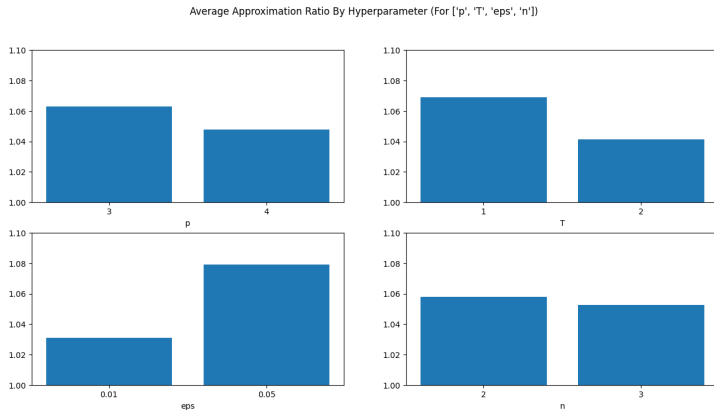


Figure: Average Approximation Ratio For p , T , ϵ , and n
Best Values are $p = 4$, $T = 2$, $\epsilon = 0.01$, and $n = 3$.

Avg. Error By Hyperparameter (2)

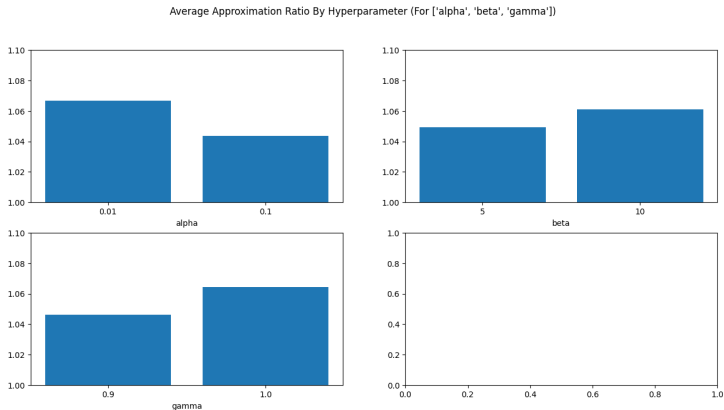


Figure: Average Approximation Ratio For α , β , and γ

Best Values are $\alpha = 0.1$, $\beta = 5$, and $\gamma = 0.9$.

Best Hyperparameters?

- Considering Hyperparameters Independently:
According to bar graphs, the best hyperparameters are:
 $p = 4$, $T = 2$, $\epsilon = 0.01$, $n = 3$, $\alpha = 0.1$, $\beta = 5$, and $\gamma = 0.9$
- Considering Hyperparameters As A Set:
 - Using the model with the above assignment of hyperparameters, the average approximation ratio was $\rho = 1$.
 - Of the 128 models used in cross-validation, 39 of them had an approximation ratio of $\rho = 1$.
- We will consider the above assignment of hyperparameters to give us our “best” model.

Q-Learning With “Best” Model Over 3 Episodes (1st Sample, 9 Vertices Each)




Figure: Q-Learning Over 3 Episodes (1st Sample, 9 Vertices Each)

Q-Learning With “Best” Model Over 3 Episodes (2nd Sample, 14 Vertices Each)

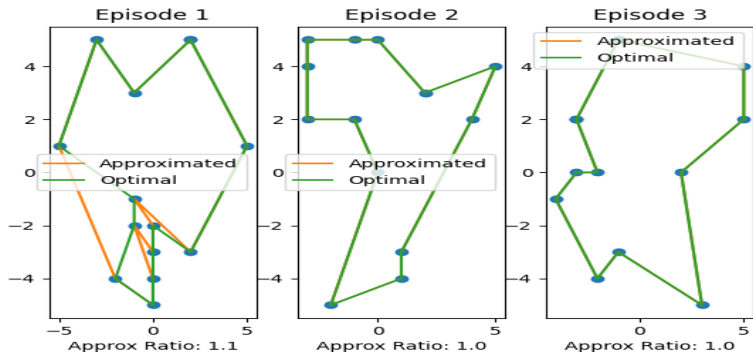


Figure: Q-Learning Over 3 Episodes (2nd Sample 14 Vertices Each)

Q-Learning With “Best” Model Over 3 Episodes (3rd Sample, 17 Vertices Each)




Figure: Q-Learning Over 3 Episodes (3rd Sample, 17 Vertices Each)



Focus Of Project:

Hanjun Dai, Elias B. Khalil, Yuyu Zhang, Bistra Dilkina, Le Song
Learning Combinatorial Optimization Algorithms over Graphs
Neural Processing Information Systems 5 April, 2017



For The Held-Karp Algorithm:

Feidiao Yang, Tiancheng Jin, Tie-Yan Liu, Xiaoming Sun, Jialing Zhang
Boosting Dynamic Programming with Neural Networks for Solving NP-hard Problems
Proceedings of Machine Learning Research 95:726-739, 2018.